

Data Cleaning and preparation

Week 4/5--Topic 4

jingyun.wang@durham.ac.uk

Handling missing data

https://pandas.pydata.org/pandas-docs/stable/user_guide/missing_data.html

Pandas refers to missing data as **NA**, which stands for not available .

A sentinel value

- For numeric data, pandas uses the floating-point value **NaN** (Not a Number) to represent missing data.
- The built-in Python **None** value is also treated as **NA** in object arrays.

```
import numpy as np
import pandas as pd
string_data = pd.Series([None, 'artichoke', np.nan, 'avocado'])
string_data.isnull()
```

```
0    True
1   False
2    True
3   False
dtype: bool
```

Table 7-1. NA handling methods

Argument	Description
<code>dropna</code>	Filter axis labels based on whether values for each label have missing data, with varying thresholds for how much missing data to tolerate.
<code>fillna</code>	Fill in missing data with some value or using an interpolation method such as <code>'ffill'</code> or <code>'bfill'</code> .
<code>isnull</code>	Return boolean values indicating which values are missing/NA.
<code>notnull</code>	Negation of <code>isnull</code> .

Filtering out missing data

On a Series, it returns the Series with only the non-null data and index values:

```
from numpy import nan as NA
data = pd.Series([1, NA, 3.5, NA, 7.9])
data.dropna() → data[data.notnull()]
```

```
0    1.0
2    3.5
4    7.9
dtype: float64
```

For dataframe, you can drop rows or columns that are all NA or only those containing any NAs. dropna by default drops any row containing a missing value:

```
data = pd.DataFrame([[1., 6.5, 3.], [1., NA, NA], [NA, NA, NA], [NA, 6.5, 3.]])
cleaned = data.dropna()
data
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

cleaned

	0	1	2
0	1.0	6.5	3.0

only want to drop rows that are all NA:

```
data.dropna(how="all")
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
3	NaN	6.5	3.0

Examples

To drop columns in the same way, pass axis=1

```
data[4] = NA  
data
```

	0	1	2	4
0	1.0	6.5	3.0	NaN
1	1.0	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	6.5	3.0	NaN

```
data.dropna(axis=1, how='all')
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

Filling in missing data

Calling `fillna()` with a constant replaces missing values with that value:

```
data.fillna(0)
```

	0	1	2	4
0	1.0	6.5	3.0	0.0
1	1.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	6.5	3.0	0.0

Calling `fillna()` with a dict, you can use a different fill value for each **column**:

```
data.fillna({1: 5, 2: 0, 4: 7})
```

	0	1	2	4
0	1.0	6.5	3.0	7.0
1	1.0	5.0	0.0	7.0
2	NaN	5.0	0.0	7.0
3	NaN	6.5	3.0	7.0

With fillna you can do lots of other things with a little creativity. For example, you might pass the **mean** or **median** value of a Series:

```
data=pd.Series([1.0,NA,3.5,NA,7])
data.fillna(data.mean())
```

```
0    1.000000
1    3.833333
2    3.500000
3    3.833333
4    7.000000
dtype: float64
```

Table 7-2. fillna function arguments

Argument	Description
value	Scalar value or dict-like object to use to fill missing values
method	Interpolation; by default 'ffill' if function called with no other arguments
axis	Axis to fill on; default axis=0
inplace	Modify the calling object without producing a copy
limit	For forward and backward filling, maximum number of consecutive periods to fill

fillna returns a new object, but you can modify the existing object in-place:

```
data.fillna(0, inplace=True)
```


Examples

The same interpolation methods available for reindexing

	0	1	2	4
0	1.0	6.5	3.0	NaN
1	1.0	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	2.5	3.5	NaN

ffill (stands for 'forward fill'): replaces the NULL values with the value from the previous row (or previous column, if the axis parameter is set to 'columns').

```
data.fillna(method="ffill")
```



```
data.fillna(method="ffill", limit=1)
```

	0	1	2	4
0	1.0	6.5	3.0	NaN
1	1.0	6.5	3.0	NaN
2	1.0	6.5	3.0	NaN
3	1.0	2.5	3.5	NaN

	0	1	2	4
0	1.0	6.5	3.0	NaN
1	1.0	6.5	3.0	NaN
2	1.0	NaN	NaN	NaN
3	NaN	2.5	3.5	NaN

Data transformation

Removing Duplicates

```
data = pd.DataFrame({'k1': ['one', 'two'] * 3 + ['two'], 'k2': [1, 1, 2, 3, 3, 4, 4]})
data
```

	k1	k2
0	one	1
1	two	1
2	one	2
3	two	3
4	one	3
5	two	4
6	two	4

The DataFrame method `data.duplicated()` returns a boolean Series indicating whether each row is a duplicate (has been observed in a previous row) or not:

```
data.duplicated()
```

```
0    False
1    False
2    False
3    False
4    False
5    False
6     True
dtype: bool
```

`drop_duplicates()` returns a DataFrame where the duplicated array is False

```
data.drop_duplicates()
```


	k1	k2
0	one	1
1	two	1
2	one	2
3	two	3
4	one	3
5	two	4

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
5	two	4	5
6	two	4	6

specify any subset of them to detect duplicates.

```
data.drop_duplicates(['k1'])
```

	k1	k2	v1
0	one	1	0
1	two	1	1



by default keep
the first
observed value
combination

```
data.drop_duplicates(["k1","k2"])
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
5	two	4	5

return the last one:

```
data.drop_duplicates(["k1","k2"], keep="last")
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
6	two	4	6

Transforming Data Using a Function or Mapping

```
data = pd.DataFrame({'food': ['bacon', 'pulled pork', 'bacon',  
                             'Pastrami', 'corned beef', 'Bacon',  
                             'pastrami', 'honey ham', 'nova lox'],  
                    'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})  
  
meat_to_animal = { 'bacon': 'pig',  
                  'pulled pork': 'pig',  
                  'pastrami': 'cow',  
                  'corned beef': 'cow',  
                  'honey ham': 'pig',  
                  'nova lox': 'salmon' }  
  
lowercased = data['food'].str.lower()  
lowercased
```

```
0      bacon  
1  pulled pork  
2      bacon  
3    pastrami  
4  corned beef  
5      bacon  
6    pastrami  
7  honey ham  
8    nova lox  
Name: food, dtype: object
```

```
data['animal'] = lowercased.map(meat_to_animal)  
data
```

	food	ounces	animal
0	bacon	4.0	pig
1	pulled pork	3.0	pig
2	bacon	12.0	pig
3	Pastrami	6.0	cow
4	corned beef	7.5	cow
5	Bacon	8.0	pig
6	pastrami	3.0	cow
7	honey ham	5.0	pig
8	nova lox	6.0	salmon

Replacing values

```
data = pd.Series([1., -999., 2., -999., -1000., 3.])  
data.replace(-999, np.nan)
```

```
0      1.0  
1      NaN  
2      2.0  
3      NaN  
4    -1000.0  
5       3.0  
dtype: float64
```

If you want to replace multiple values at once, you instead pass a list and then the substitute value:

```
data.replace([-999, -1000], np.nan)
```

To use a different replacement for each value, pass a list of substitutes:

```
data.replace([-999, -1000], [np.nan, 0])
```

The argument passed can also be a dict:

```
data.replace({-999: np.nan, -1000: 0})
```

Renaming Axis Indexes

```
data = pd.DataFrame(np.arange(12).reshape((3, 4)),
                    index=['Ohio', 'Colorado', 'New York'],
                    columns=['one', 'two', 'three', 'four'])
##Like a Series, the axis indexes have a map method
transform = lambda x: x[:4].upper()
data.index.map(transform)
```

```
Index(['OHIO', 'COLO', 'NEW '], dtype='object')
```

```
##You can assign to index, modifying the DataFrame in-place:
data.index = data.index.map(transform)
data
```

	one	two	three	four
OHIO	0	1	2	3
COLO	4	5	6	7
NEW	8	9	10	11

##If you want to create a transformed version of a dataset without modifying the original, a useful method is rename :

```
data.rename(index=str.title, columns=str.upper)
```

	ONE	TWO	THREE	FOUR
Ohio	0	1	2	3
Colo	4	5	6	7
New	8	9	10	11

##Notably, rename can be used in conjunction with a dict-like object providing new values for a subset of the axis labels:

```
data.rename(index={'OHIO': 'INDIANA'},  
            columns={'three': 'peekaboo'})
```

	one	two	peekaboo	four
INDIANA	0	1	2	3
COLO	4	5	6	7
NEW	8	9	10	11

#Should you wish to modify a dataset in-place, :

```
data.rename(index={'OHIO': 'INDIANA'}, inplace=True)  
data
```

	one	two	three	four
INDIANA	0	1	2	3
COLO	4	5	6	7
NEW	8	9	10	11

Discretization and binning

Continuous data is often discretized or otherwise separated into “bins” for analysis.

Suppose you have data about a group of people in a study, and you want to group them into discrete age buckets:

```
import pandas as pd
ages=[20,22,25,27,21,23,37,31,61,45,41,32]
bins=[18,25,35,60,100]
cats=pd.cut(ages,bins)      special Categorical object
cats
```

```
[(18, 25], (18, 25], (18, 25], (25, 35], (18, 25], ..., (25, 35], (60, 100], (35, 60], (35, 60], (25, 35]]
Length: 12
Categories (4, interval[int64]): [(18, 25] < (25, 35] < (35, 60] < (60, 100]]
```

You can change which side is closed by passing **right=False** :

```
pd.cut(ages, [18, 26, 36, 61, 100], right=False) Indicates whether bins includes the rightmost edge or not.
```

```
[[18, 26), [18, 26), [18, 26), [26, 36), [18, 26), ..., [26, 36), [61, 100), [36, 61), [36, 61), [26, 36)]
Length: 12
Categories (4, interval[int64]): [[18, 26) < [26, 36) < [36, 61) < [61, 100))
```

<https://pandas.pydata.org/docs/reference/api/pandas.Categorical.html>

```
type(cats)
```

```
pandas.core.arrays.categorical.Categorical
```

```
cats.categories
```

```
IntervalIndex([(18, 25], (25, 35], (35, 60], (60, 100]]  
              closed='right',  
              dtype='interval[int64]')
```

```
cats.codes
```

```
array([0, 0, 0, 1, 0, 0, 2, 1, 3, 2, 2, 1], dtype=int8)
```

```
pd.value_counts(cats)
```

```
(18, 25]      5  
(35, 60]      3  
(25, 35]      3  
(60, 100]     1  
dtype: int64
```

Attributes

`categories`

The categories of this categorical.

`codes`

The category codes of this categorical.

`ordered`

Whether the categories have an ordered relationship.

`dtype`

The `CategoricalDtype` for this instance.

You can also pass your own bin names by passing a list or array to the labels options:

```
group_names = ['Youth', 'YoungAdult', 'MiddleAged', 'Senior']  
pd.cut(ages, bins, labels=group_names)
```

```
[Youth, Youth, Youth, YoungAdult, Youth, ..., YoungAdult, Senior, MiddleAged, MiddleAged, YoungAdult]  
Length: 12  
Categories (4, object): [Youth < YoungAdult < MiddleAged < Senior]
```

Detecting and Filtering Outliers

Filtering or transforming outliers is largely a matter of applying array operations. Consider a DataFrame with some normally distributed data:

```
data = pd.DataFrame(np.random.randn(1000, 4))
data.describe()
```

	0	1	2	3
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.004846	-0.021556	0.015159	0.039501
std	1.015053	0.981478	1.020946	1.008537
min	-3.259632	-2.760693	-2.890439	-3.276493
25%	-0.655477	-0.699397	-0.704225	-0.666109
50%	-0.022993	-0.016327	0.003854	0.087594
75%	0.668354	0.615626	0.719253	0.708329
max	3.393005	2.489701	3.240661	3.325200

find values in one of the columns exceeding 3 in absolute value:

```
col=data[2]
col[np.abs(col)>3]
```

```
675    3.099820
823    3.240661
Name: 2, dtype: float64
```

To select all rows having a value exceeding 3 or – 3, you can use the boolean DataFrame:

```
data[(np.abs(data) > 3).any(1)]
```

	0	1	2	3
228	3.189155	-0.937648	-0.634416	-0.727816
437	-3.145205	-1.085325	-0.787895	-0.482588
440	-3.171539	-0.930211	0.108549	0.362863
488	-0.255291	-1.100287	-0.940546	-3.223819
556	1.089107	0.094619	0.342310	-3.276493
594	3.158020	1.849592	-1.689944	-1.264951
668	-0.759775	0.152895	1.591284	3.131849
675	-0.947180	-0.175683	3.099820	-0.652692
694	1.495299	-1.245924	-1.105193	3.325200
724	3.393005	-0.527329	-0.014588	1.153746
760	3.354250	-1.817975	0.163152	0.820375
823	-0.290860	0.862962	3.240661	0.632690
972	-3.259632	-0.901415	1.453389	1.262064

Values can be set based on these criteria. Here is code to cap values outside the interval -3 to 3 :

```
data[np.abs(data) > 3] = np.sign(data) * 3
data.describe()
```

	0	1	2	3
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.004328	-0.021556	0.014818	0.039544
std	1.009873	0.981478	1.019916	1.005556
min	-3.000000	-2.760693	-2.890439	-3.000000
25%	-0.655477	-0.699397	-0.704225	-0.666109
50%	-0.022993	-0.016327	0.003854	0.087594
75%	0.668354	0.615626	0.719253	0.708329
max	3.000000	2.489701	3.000000	3.000000

The statement `np.sign(data)` produces 1 and -1 values based on whether the values in data are positive or negative

Permutation and random sampling

`DataFrame.take(indices, axis=0, **kwargs)`

Return the elements in the given positional indices along an axis.

```
df = pd.DataFrame(np.arange(5 * 4).reshape((5, 4)))  
df
```

	0	1	2	3
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11
3	12	13	14	15
4	16	17	18	19

```
sampler = np.random.permutation(5)  
sampler
```

```
array([3, 0, 4, 1, 2])
```

```
df.take(sampler)
```

	0	1	2	3
3	12	13	14	15
0	0	1	2	3
4	16	17	18	19
1	4	5	6	7
2	8	9	10	11

To select a random subset without replacement, you can use the **sample** method on Series and DataFrame:

```
df.sample(n=3)
```

	0	1	2	3
1	4	5	6	7
4	16	17	18	19
3	12	13	14	15

To generate a sample with replacement (to allow repeat choices), pass to sample :

```
choices = pd.Series([5, 7, -1, 6, 4])
draws = choices.sample(n=10, replace=True)
draws
```

```
4      4
3      6
0      5
0      5
2     -1
2     -1
3      6
1      7
0      5
1      7
dtype: int64
```




Durham
University

String manipulation

String Object Methods

Table 7-3. Python built-in string methods

Method	Description
<code>count</code>	Return the number of non-overlapping occurrences of substring in the string.
<code>endswith</code>	Returns <code>True</code> if string ends with suffix.
<code>startswith</code>	Returns <code>True</code> if string starts with prefix.
<code>join</code>	Use string as delimiter for concatenating a sequence of other strings.
<code>index</code>	Return position of first character in substring if found in the string; raises <code>ValueError</code> if not found.
<code>find</code>	Return position of first character of <i>first</i> occurrence of substring in the string; like <code>index</code> , but returns <code>-1</code> if not found.
<code>rfind</code>	Return position of first character of <i>last</i> occurrence of substring in the string; returns <code>-1</code> if not found.
<code>replace</code>	Replace occurrences of string with another string.
<code>strip</code> , <code>rstrip</code> , <code>lstrip</code>	Trim whitespace, including newlines; equivalent to <code>x.strip()</code> (and <code>rstrip</code> , <code>lstrip</code> , respectively) for each element.
<code>split</code>	Break string into list of substrings using passed delimiter.
<code>lower</code>	Convert alphabet characters to lowercase.
<code>upper</code>	Convert alphabet characters to uppercase.
<code>casefold</code>	Convert characters to lowercase, and convert any region-specific variable character combinations to a common comparable form.
<code>ljust</code> , <code>rjust</code>	Left justify or right justify, respectively; pad opposite side of string with spaces (or some other fill character) to return a string with a minimum width.

Regular Expressions

Table 7-4. Regular expression methods

Method	Description
<code>findall</code>	Return all non-overlapping matching patterns in a string as a list
<code>finditer</code>	Like <code>findall</code> , but returns an iterator
<code>match</code>	Match pattern at start of string and optionally segment pattern components into groups; if the pattern matches, returns a match object, and otherwise <code>None</code>
<code>search</code>	Scan string for match to pattern; returning a match object if so; unlike <code>match</code> , the match can be anywhere in the string as opposed to only at the beginning
<code>split</code>	Break string into pieces at each occurrence of pattern
<code>sub</code> , <code>subn</code>	Replace all (<code>sub</code>) or first <code>n</code> occurrences (<code>subn</code>) of pattern in string with replacement expression; use symbols <code>\1</code> , <code>\2</code> , ... to refer to match group elements in the replacement string

Vectorized String Functions in pandas

Table 7-5. Partial listing of vectorized string methods

Method	Description
<code>cat</code>	Concatenate strings element-wise with optional delimiter
<code>contains</code>	Return boolean array if each string contains pattern/regex
<code>count</code>	Count occurrences of pattern
<code>extract</code>	Use a regular expression with groups to extract one or more strings from a Series of strings; the result will be a DataFrame with one column per group
<code>endswith</code>	Equivalent to <code>x.endswith(pattern)</code> for each element
<code>startswith</code>	Equivalent to <code>x.startswith(pattern)</code> for each element
<code>findall</code>	Compute list of all occurrences of pattern/regex for each string
<code>get</code>	Index into each element (retrieve <i>i</i> -th element)
<code>isalnum</code>	Equivalent to built-in <code>str.isalnum</code>
<code>isalpha</code>	Equivalent to built-in <code>str.isalpha</code>
<code>isdecimal</code>	Equivalent to built-in <code>str.isdecimal</code>
<code>isdigit</code>	Equivalent to built-in <code>str.isdigit</code>
<code>islower</code>	Equivalent to built-in <code>str.islower</code>
<code>isnumeric</code>	Equivalent to built-in <code>str.isnumeric</code>
<code>isupper</code>	Equivalent to built-in <code>str.isupper</code>
<code>join</code>	Join strings in each element of the Series with passed separator
<code>len</code>	Compute length of each string
<code>lower</code> , <code>upper</code>	Convert cases; equivalent to <code>x.lower()</code> or <code>x.upper()</code> for each element

Method	Description
<code>match</code>	Use <code>re.match</code> with the passed regular expression on each element, returning <code>True</code> or <code>False</code> whether it matches.
<code>extract</code>	Extract captured group element (if any) by index from each string
<code>pad</code>	Add whitespace to left, right, or both sides of strings
<code>center</code>	Equivalent to <code>pad(side='both')</code>
<code>repeat</code>	Duplicate values (e.g., <code>s.str.repeat(3)</code> is equivalent to <code>x * 3</code> for each string)
<code>replace</code>	Replace occurrences of pattern/regex with some other string
<code>slice</code>	Slice each string in the Series
<code>split</code>	Split strings on delimiter or regular expression
<code>strip</code>	Trim whitespace from both sides, including newlines
<code>rstrip</code>	Trim whitespace on right side
<code>lstrip</code>	Trim whitespace on left side

A data cleaning example

<https://www.kaggle.com/parulpandey/2020-it-salary-survey-for-eu-region>

IT Salary Survey EU 2020

In order to Explore the situation of IT salary, clean the dataset

To explore the following questions :

- (1) Which are the top 20 positions among the respondents of the surveys?
(Compare the results in 2019 and 2020.)
- (2) Which are the top 10 main programming languages among the respondents? (Compare the results in 2019 and 2020.)
- (3) What affect the annual income of the respondents

What happen if there are more than one that take the 10th place ?

Typical data cleaning in Natural Language Processing

<https://www.machinelearningplus.com/nlp/natural-language-processing-guide/>

More than 80% of the data available today is **Unstructured Data**.

The **texts, videos, images** which cannot be represented in a tabular form (or in any consistent structured data model) constitute unstructured Data.

Natural Language Toolkit (NLTK)

Text Pre-processing in NLP--nlTK and spacy

The raw text data often referred to as text corpus has a lot of noise. There are **punctuation**, **suffices** and **stop words** that do not give us any information. Text Processing involves preparing the text corpus to make it more usable for NLP tasks.

For pre-processing of the raw text, follow the steps

1. Tokenization

The words of a text document/file separated by spaces and punctuation are called as tokens. The process of extracting tokens from a text file/document is referred as tokenization.

2. Sanitization of the strings - making text lowercase, removing all punctuation, replacing all whitespace/newlines with single spaces (' ')

3. Removal of **stopwords** (e.g. "and", "the")

5. Stemming --reducing a word to its '**root form**'.

6. Lemmatization - reducing each word to a **root/base form** (*similar to stemming, except that the root word is correct and always meaningful.*)